

INTELLIGENT TRAFFIC MANAGEMENT SYSTEM

ITMS

System & Technical Documentation

-49%Emergency Travel
Time**0.911**ML R² (30s)**132**

Tests Passing

7/7

Phases Complete

Team: Bitcoders**Event:** SISTec Innovation Hackathon 2026**Problem Statement:** IS-8**Date:** September 27, 2026

Documentation generated from source-code analysis
All metrics measured from actual SUMO simulation runs — nothing fabricated



ITMS — System & Technical Documentation

Intelligent Traffic Management System *Predictive Rolling Green Corridor for Emergency Vehicles*

Team	Bitcoders
Event	SISTec Innovation Hackathon 2026
Problem Statement	IS-8
Documentation Date	September 27, 2026
Implementation Status	All 7 Phases Complete

Table of Contents

- [Executive Summary](#)
- [Problem Statement](#)
- [System Objectives](#)
- [System Architecture](#)
- [Technology Stack](#)
- [Simulation Architecture \(SUMO + TraCI\)](#)
- [Traffic Intelligence Layer](#)
- [Emergency Vehicle Management](#)
- [Dynamic Routing \(A* Algorithm\)](#)
- [AI Traffic Prediction \(XGBoost\)](#)
- [Predictive Green Corridor](#)
- [Safety Constraint Engine](#)
- [Closed-Loop Optimization](#)

- [Database Architecture](#)
- [API Reference](#)
- [Real-Time Communication \(WebSocket\)](#)
- [Testing & Verification](#)
- [Performance Metrics](#)
- [Deployment Guide](#)
- [Limitations & Known Issues](#)
- [Implementation Status Matrix](#)

1. Executive Summary

The **Intelligent Traffic Management System (ITMS)** is a closed-loop, AI-augmented traffic orchestration platform. Its core innovation is the **Predictive Rolling Green Corridor** — a system that pre-clears traffic signals ahead of an emergency vehicle by predicting where the vehicle will be in the future, rather than simply reacting when the vehicle arrives at a signal.

ITMS operates on a digital traffic twin built with **Eclipse SUMO** (Simulation of Urban MObility), controlled in real time via the **TraCI** protocol. A **Node.js + Fastify** backend orchestrates routing, prediction, and corridor planning. An **XGBoost-based Python microservice** provides 30–120 second traffic volume forecasts. A **Next.js command center** provides operators with a live mission-control interface.

Verified Performance Results (from simulation)

Metric	Baseline (No ITMS)	ITMS	Improvement
Emergency Travel Time	162 s	82 s	−49.4%
Average Vehicle Delay	21.94 s	15.66 s	−28.6%
Average Queue Length	25.59	21.78	−14.9%
Average Speed	6.09 m/s	6.36 m/s	+4.4%
Throughput	18.65 /h	31.86 /h	+70.8%

All values are measured from actual SUMO simulation runs. None are fabricated.

2. Problem Statement

Emergency vehicles — ambulances, fire engines, police — lose critical minutes to urban traffic congestion. Every 10-second delay in ambulance response increases cardiac arrest mortality by approximately 10%.

Traditional traffic signal preemption is **reactive**: it only changes a signal when the vehicle physically arrives at an intersection, creating:


- **Shockwave congestion** — sudden green signals cause downstream queuing.
- **Secondary queue buildup** — cross-traffic blocked without warning.
- **Pedestrian hazards** — abrupt signal changes endanger pedestrians.
- **Corridor gaps** — vehicles must stop at each signal, defeating the purpose.


ITMS solves this by being **predictive**: it plans signal states *before* the emergency vehicle arrives, coordinating an entire corridor of intersections simultaneously.

3. System Objectives

- **Digital Twin**: Create an authoritative real-time simulation of city traffic using SUMO.
- **Traffic Intelligence**: Continuously measure vehicle counts, speeds, queues, and congestion across all road segments.
- **Predictive ML**: Forecast traffic conditions at each intersection up to 120 seconds in the future.
- **Dynamic Routing**: Route emergency vehicles using A* search over live, congestion-weighted road graphs.
- **Green Corridor**: Pre-clear intersections ahead of the emergency vehicle based on predicted ETAs.
- **Safety**: Enforce explicit safety constraints — no conflicting movements, bounded red/green extensions, clearance intervals.
- **Closed Loop**: Continuously re-evaluate routes and corridors as conditions change.
- **Command Center**: Provide operators with a live, data-driven mission control interface.

4. System Architecture





Key Architectural Principles

- **Simulation-first:** SUMO is the authoritative source of truth. No values are fabricated.
- **ML for prediction, algorithms for decisions:** XGBoost predicts traffic; A* and corridor logic make decisions.
- **Safe degradation:** If ML service, database, or SUMO is unavailable, the system falls back gracefully — never crashing.
- **Closed-loop:** Every re-evaluation loop runs on sim-time cadence, serialized with the corridor executor.

5. Technology Stack

Layer	Technology	Version	Role	Status
Frontend	Next.js	15.5.6	Command Center UI	IMPLEMENTED
Frontend	React	19	UI components	IMPLEMENTED
Frontend	Tailwind CSS v4	4.x	Styling	IMPLEMENTED
Frontend	Framer Motion	—	Animations	IMPLEMENTED
Frontend	Recharts	—	Analytics charts	IMPLEMENTED
Backend	Node.js	24.7	Runtime	IMPLEMENTED
Backend	Fastify	—	HTTP framework	IMPLEMENTED
Backend	TypeScript	—	Type safety	IMPLEMENTED
Backend	@fastify/cors	—	CORS support	IMPLEMENTED
Backend	@fastify/websocket	—	WS support	IMPLEMENTED
Database	PostgreSQL	16	Persistent storage	IMPLEMENTED
Database	PostGIS	3.4	Spatial extensions	IMPLEMENTED
Simulation	Eclipse SUMO	1.27.1	Traffic simulation	IMPLEMENTED
Simulation	TraCI (TCP)	—	SUMO control API	IMPLEMENTED
ML	Python	3.11	ML runtime	IMPLEMENTED
ML	XGBoost	—	Traffic prediction	IMPLEMENTED
ML	FastAPI	—	Prediction HTTP API	IMPLEMENTED
ML	scikit-learn	—	Preprocessing	IMPLEMENTED

ML	Pandas / NumPy	—	Data processing	IMPLEMENTED
Routing	A* Algorithm	—	Emergency routing	IMPLEMENTED
AI	Vultr Serverless	—	AI Copilot (LLM)	IMPLEMENTED
Visualization	SVG (custom)	—	SUMO network render	IMPLEMENTED

Note: MapLibre GL was removed. Google Maps was integrated then also removed. The final map is a pure SVG renderer of the actual SUMO network geometry, driven by real TraCI coordinates.

6. Simulation Architecture

SUMO Network

The system supports two network profiles:

Grid Profile (Default)

A synthetic 3×2 signalized intersection grid with 6 controlled junctions (I1–I6) plus boundary stubs (W1, W2, E1, E2, S1–S3, N1–N3). Built from

`itms.nod.xml` + `itms.edg.xml` via `netconvert`.

City Profile (Bhopal)

A real OpenStreetMap-derived network of central Bhopal, georeferenced with UTM projection. Generated via `npm run fetch:city && npm run build:city-network`. Contains 883+ road segments and 21+ traffic light systems.

Simulation Scenarios

Scenario ID	Description	File
<code>baseline</code>	Normal traffic, no emergency	<code>scenarios/baseline/baseline.sumocfg</code>

<code>emergency</code>	Normal traffic, emergency vehicles dynamically spawned	<code>scenarios/emergency/emergency.sumocfg</code>
<code>emergency_low</code>	Low demand traffic	<code>scenarios/emergency/emergency-low.sumocfg</code>
<code>emergency_high</code>	High demand traffic	<code>scenarios/emergency/emergency-high.sumocfg</code>

Emergency Fleet Definition

Defined in `simulation/sumo/network/emergency-fleet.add.xml` :

- `emergency.ambulance` — vClass: `emergency`
- `emergency.fireengine` — vClass: `emergency`
- `emergency.police` — vClass: `authority`

TraCI Integration

The backend uses a **raw TCP TraCI client** (no third-party library) implementing the TraCI binary protocol verified against SUMO 1.27.1.

TraCI Variables Read Per Step:

- `EDGE_VEHICLE_COUNT` — vehicle count per edge
- `EDGE_MEAN_SPEED` — average speed per edge
- `EDGE_HALTING_NUMBER` — halted vehicles per edge
- `EDGE_OCCUPANCY` — occupancy ratio per edge
- `TL_CURRENT_PHASE` — current signal phase
- `TL_CURRENT_PROGRAM` — current program ID
- `TL_PHASE_DURATION` — phase duration
- `TL_NEXT_SWITCH` — time to next phase switch
- `TL_RED_YELLOW_GREEN_STATE` — full RYG state string
- `VAR_POSITION` (vehicles) — x/y coordinates
- `VAR_SPEED` (vehicles) — speed in m/s
- `VAR_ANGLE` (0x43) — vehicle heading angle

- `VAR_ROAD_ID` (vehicles) — current edge
- `VAR_LANEPOSITION` (vehicles) — position on lane
- `VAR_TIMELOSS` (vehicles) — time lost to congestion

TraCI Commands Sent:

- `setRedYellowGreenState` — set signal state string
- `setProgram` — restore normal signal program
- `setSpeed` — control vehicle speed (emergency spawning)
- `add` / `route.add` — spawn emergency vehicles

7. Traffic Intelligence Layer

The traffic collector runs **lock-step** with the simulation: after each SUMO step, it reads the complete traffic state before the next step begins.

Metrics Computed Per Segment

- **Vehicle Count** — raw SUMO edge domain count
- **Average Speed** — SUMO `meanSpeed`, gated on `vehicleCount > 0`
- **Halting Count** — SUMO `haltingNumber`
- **Occupancy** — SUMO `occupancy` ratio [0..1]
- **Queue Length** — `haltingNumber` mapped to queue estimate
- **Flow Rate** — derived from observed vehicle edge transitions between ticks
- **Congestion Level** — deterministic classification:

```
LOW:      occupancy < 0.15 OR speed_ratio > 0.6 AND queue < 2

MEDIUM:   occupancy < 0.30 OR speed_ratio > 0.35 AND queue < 6

HIGH:      occupancy < 0.45 OR speed_ratio > 0.15 AND queue < 12

CRITICAL: otherwise
```

Thresholds are configurable via `TRAFFIC_CONGESTION_` environment variables.*

Persistence & Broadcasting

- Snapshots persisted to `traffic_snapshots` and `signal_snapshots` every `TRAFFIC_PERSIST_EVERY_TICKS` (default: 5 ticks).
- WebSocket broadcast paced at `TRAFFIC_EVENT_INTERVAL_MS` (default: 200ms / ~5 Hz).
- Run identity tracked — all snapshots linked to a `simulation_runs` row.

8. Emergency Vehicle Management

Emergency vehicles are created **dynamically** via API. There are no hardcoded ambulances in the scenarios.

Creation Workflow

```
POST /api/emergency
```

↓

```
Validate type, priority, origin, destination junctions
```

↓

```
Compute A* route over live traffic graph
```

↓

```
Persist emergency_vehicle + emergency_event + route + route_segments
```

↓

```
TraCI: add route to SUMO
```

↓

TraCI: add vehicle to SUMO

↓

Per-step tracking: position, speed, angle, route_index

↓

Per-step ETA calculation for each upcoming junction

↓

Arrival detection via SUMO arrived-vehicle list (VAR_ARRIVED_VEHICLES_I

Supported Types

Type	SUMO vClass	Priority Levels
ambulance	emergency	critical, high, normal
fire_engine	emergency	critical, high, normal
police	authority	critical, high, normal

ETA Calculation

$$ETA = \sum (\text{segment_length} / \text{effective_speed}) \times \text{progress_factor}$$

$$\text{effective_speed} = \text{clamp}(\text{measured_avg_speed}, \text{free_flow} \times 0.1, \text{free_flow})$$

```
progress_factor = clamp(demonstrated_pace / expected_pace, 0.5, 1.0)
```

9. Dynamic Routing

Road Graph

```
Nodes = Intersections (junction IDs)

Edges = Directed road segments (SUMO edge IDs)
```

Edge Cost Model

```
cost = (segment_length / effective_speed)

× congestion_adjustment_factor
```

Where `congestion_adjustment_factor` accounts for measured occupancy and speed ratio on that segment. No congestion data → uses free-flow speed.

A* Implementation

- **Heuristic:** Euclidean distance / max_speed (admissible — never overestimates)
- **Priority queue:** Binary heap
- **Tie-breaking:** Deterministic by junction ID
- **All 30 junction pairs** of the 3×2 grid are verified routable

Route Switching (Closed Loop)

The closed-loop service re-evaluates routes on a sim-time cadence with hysteresis safeguards:

Safeguard	Value
-----------	-------

Minimum improvement (absolute)	8 seconds
Minimum improvement (relative)	15% of current ETA
Cooldown between switches	30 seconds
Maximum switches per emergency	3
Near-destination guard	No switch within 15s of destination
Crossing-intersection guard	No switch while routeIndex < 0

Every route switch is persisted to `emergency_route_switches` and broadcast as `route:switched`.

10. AI Traffic Prediction

Dataset

Property	Value
Total Rows	36,000
Total SUMO Runs	40
Collection Interval	2 seconds
Split Method	Run-level 70/15/15 (seeded, no temporal leakage)
Train Runs	28 runs
Validation Runs	6 runs
Test Runs	6 runs (gen01, gen07, gen08, gen14, gen15, gen17)

Features

Feature	Description
<code>vehicle_count</code>	Current junction approach count
<code>avg_speed_mps</code>	Average approach speed
<code>queue_length</code>	Halting vehicle count
<code>density</code>	Vehicles per meter

<code>flow_rate_per_hour</code>	Edge departure turnover
<code>signal_phase</code>	Current signal phase index
<code>signal_green_fraction</code>	Green ratio in current program
<code>cycle_position_s</code>	Position in signal cycle
<code>hour</code>	Simulated hour of day
<code>day_of_week</code>	Simulated day of week
<code>count_lag_</code> / <code>queue_lag_</code>	Lagged traffic features

Model Architecture

One **XGBRegressor** per prediction horizon (4 models total):

Horizon	Target
+30s	<code>target_count_30</code>
+60s	<code>target_count_60</code>
+90s	<code>target_count_90</code>
+120s	<code>target_count_120</code>

Model Parameters

```
{
  "objective": "reg:squarederror",

  "n_estimators": 600,

  "max_depth": 8,

  "learning_rate": 0.06,

  "subsample": 0.9,
```

```
"colsample_bytree": 0.9,  
  
"reg_lambda": 1.0,  
  
"random_state": 42,  
  
"n_jobs": 4,  
  
"tree_method": "hist"  
  
}
```

Measured Test Metrics (Held-Out Runs)

Horizon	MAE	RMSE	R ²
+30s	3.580	5.159	0.911
+60s	4.605	6.881	0.873
+90s	4.969	8.343	0.846
+120s	6.942	11.830	0.739

Source: `services/prediction/evaluation/evaluation_report.json` —
generated 2026-09-23

Inference Performance

Metric	Value
Single horizon latency	1.01 ms
All four horizons latency	4.53 ms
Measured calls	2,000

Fallback Behavior (Rule 9)

If the ML service is unavailable:

- **Recent valid prediction** — use the last valid ML prediction if not stale
- **Deterministic trend** — extrapolate the measured 10s count trend with distance-damped growth

Fallback source is always labelled `source: "fallback"` in the API response.

11. Predictive Green Corridor

The Predictive Rolling Green Corridor is the **core USP** of ITMS.

How It Works

```
Emergency Route Computed (A*)
```

```
↓
```

```
Extract Controlled Intersections on Route
```

```
↓
```

```
For Each Intersection:
```

```
├─ Compute ETA from live position + speed
```

```
├─ Get live signal state (phase, next switch)
```

```
├─ Get ML traffic prediction at ETA
```

```
├─ Determine Window Mode:
```

```
├─ SWITCH - approach is red → give green at [ETA-lead, ETA+trail]
```

└─ EXTEND — approach is green but will switch too early → extend

└─ NOOP — normal green covers the ETA → no command needed

└─ PENDING — ETA beyond horizon → plan later (rolling)

↓

Safety Validation (8 checks)

↓

Corridor ACTIVE → Apply Signal States via TraCI

↓

Per-Step Rolling Executor:

└─ APPLY phase → set TraCI signal state

└─ CLEARANCE → insert yellow if cutting live green

└─ RESTORE → restore normal program on passage

└─ REPLAN → re-plan PENDING windows as vehicle advances

↓

Corridor COMPLETED when vehicle arrives

Window Timing Parameters

Parameter	Default	Environment Variable
Green lead time	12 s	<code>CORRIDOR_GREEN_LEAD_S</code>
Green trail time	12 s	<code>CORRIDOR_GREEN_TRAIL_S</code>
Min green window	8 s	<code>CORRIDOR_MIN_GREEN_WINDOW_S</code>
Max green window	30 s	<code>CORRIDOR_MAX_GREEN_WINDOW_S</code>
Max green extension	20 s	<code>CORRIDOR_MAX_GREEN_EXTENSION_S</code>
Max red extension	45 s	<code>CORRIDOR_MAX_RED_EXTENSION_S</code>
Yellow clearance	3 s	<code>CORRIDOR_CLEARANCE_YELLOW_S</code>
ETA plan horizon	60 s	<code>CORRIDOR_ETA_PLAN_HORIZON_S</code>
Replan threshold	4 s	<code>CORRIDOR_REPLAN_ETA_THRESHOLD_S</code>
Downstream limit	0.85	<code>CORRIDOR_DOWNSTREAM_OCCUPANCY_LIMIT</code>

Corridor Signal States

Corridor green states are **conflict-free by construction**:

- Green is given **only** to the emergency approach's link indices (from `net.xml` mapping)
- **All other movements are red**
- This is re-validated by the safety engine before execution

12. Safety Constraint Engine

8 explicit safety checks are enforced before any corridor signal command:

Check	Description
Priority Gate	Emergency priority must be $\geq$ <code>corridorMinPriority</code>

Window Bounds	Window must be within [min, max] duration
Max Red Extension	Cross-traffic cannot be held red longer than <code>corridorMaxRedExtensionS</code>
Yellow Clearance	If cutting a live green, a yellow phase is inserted
Conflict-Free State	Corridor state must not give green to conflicting movements
Downstream Capacity	If downstream occupancy > 0.85, corridor may be deferred
Corridor Conflict	If junction already commanded by higher-priority corridor, yield
Bounded Override	Corridor must not override the normal program indefinitely

If validation fails: fallback to safe signal strategy and log the rejection reason.

13. Closed-Loop Optimization

The closed-loop service runs a **serialized per-step control loop** on sim-time cadence:

```

Every loopEvalIntervals (default: 3s sim-time):
├─ Refresh ML predictions for all junctions

├─ For each active emergency:

├─ Update live position + speed + route_index

├─ Recalculate ETAs to all upcoming junctions

├─ Every routeReevalIntervals (default: 5s):
|   └─ Run A* with current traffic weights

```

```

|   └─ Check hysteresis safeguards

|   └─ If better route found → switch route

└─ If corridor needs replan → trigger corridor replan

```

Multi-Emergency Priority Policy

Scenario	Behavior
Junction actively commanded by Corridor A	Corridor B cannot take it
Junction pending for Corridor B (higher priority)	Corridor B wins
Junction pending for Corridor B (lower/equal priority)	Corridor A keeps it
Equal priority	First-come (determined by corridor activation order)

Baseline vs ITMS Comparison

```
POST /api/scenarios/compare
```

```
↓
```

```
Sequential Run 1: BASELINE
```

- Same deterministic SUMO seed
- Same warmup duration
- Emergency created at same sim time
- Normal routing + normal signal timing

- No ITMS features

↓

Sequential Run 2: ITMS

- Identical initial conditions
- A* dynamic routing + corridor + closed loop

↓

Delta computation: time, delay, queue, speed, throughput

Persisted to comparisons table

Both runs are reproducible — rerunning produces identical values.

14. Database Architecture

Migration History

Migration	Phase	Tables Created
001_phase2.sql	Traffic Intelligence	simulation_runs , intersections , roads , road_segments , traffic_signals , signal_phases , vehicles , traffic_snapshots , signal_snapshots
002_phase3.sql	Emergency + Routing	emergency_vehicles , emergency_events , routes , route_segments
003_phase4.sql	ML Predictions	traffic_predictions

004_phase5.sql	Green Corridor	green_corridors , corridor_signals
005_phase6.sql	Closed Loop	simulation_metrics , emergency_route_switches
006_phase7.sql	Command Center	comparisons
007_mobile_integration.sql	Mobile API	drivers , fleet_vehicles , driver_vehicle_assignments , hospitals , police_zones , emergency_evidence , emergency_verifications , ai_verification_results , manual_verification_decisions , audit_events , driver_telemetry
008_roadside_devices.sql	IoT Devices	Roadside device tables
009_emergency_addresses.sql	Addresses	Address fields
010_fix_driver_passwords.sql	Hotfix	Password hash fix

Complete Entity List

Table	Purpose	Key Columns
simulation_runs	Track each SUMO simulation run	id , scenario , status , sumo_version
intersections	Static junction registry	id (SUMO junction id), kind , controlled , geom (PostGIS Point)
roads	Undirected road pairs	id , from_junction , to_junction
road_segments	Directed SUMO edges	id , lane_count , length_m , max_speed_mps , geom (PostGIS LineString)
traffic_signals	Signal controllers	id , intersection_id , program_id

signal_phases	Static signal phase definitions	signal_id , phase_index , duration_s , state
vehicles	Per-run vehicle registry	run_id , vehicle_id , edge_id , position_x/y , speed_mps
traffic_snapshots	Per-step traffic measurements	run_id , segment_id , sim_time_s , vehicle_count , congestion
signal_snapshots	Per-step signal states	run_id , signal_id , phase_index , state , queue_length
emergency_vehicles	Emergency vehicle entities	vehicle_id , type , priority , status , origin/dest junction
emergency_events	Dispatch events	vehicle_id , route_id , status , created/activated/arrived_at
routes	Computed A* routes	algorithm , edge_count , total_length_m , estimated_travel_time_s
route_segments	Ordered route edges	route_id , sequence_index , segment_id , cost_seconds
traffic_predictions	ML prediction results	run_id , junction_id , horizon_s , predicted_vehicle_count , source
green_corridors	Corridor lifecycle	event_id , status , junction_count , planned/activated/completed_at
corridor_signals	Per-junction corridor plans	corridor_id , signal_id , eta_seconds , planned_green_start/end_s , mode , status
simulation_metrics	Per-run performance measurements	run_id , mode , emergency_travel_time_s

		avg_vehicle_delay_s , throughput_per_hour
emergency_route_switches	Route switch audit log	event_id , from/to_route_id , old/new_eta_s , reason
comparisons	Baseline vs ITMS comparison results	job_id , delta_travel_time_s , delta_avg_delay_s
drivers	Mobile app driver accounts	driver_code , name , email , role , status
fleet_vehicles	Physical emergency vehicles	vehicle_code , registration_number , vehicle_type , status
hospitals	Hospital registry	name , code , latitude/longitude , nearest_junction_id , available_beds
emergency_verifications	Photo verification requests	request_id , event_id , status , is_corridor_authorized
ai_verification_results	AI image analysis verdicts	verdict , confidence_score , detected_features
audit_events	Full system audit log	action , actor_id , actor_type , details

Spatial Data

All geometry uses **SRID 0** (undefined CRS) — SUMO coordinates are not geo-referenced. The PostGIS extension is used for spatial indexing and geometry types but not for geographic projection operations.

15. API Reference

All endpoints are on the Fastify server at <http://127.0.0.1:3000>.

System Endpoints

Method	Path	Purpose
GET	/health	Backend health probe
GET	/api/network/geometry	SUMO network geometry (SVG rendering data)
GET	/api/system	Overall system status
GET	/api/decisions	AI decision event log

Simulation Endpoints

Method	Path	Purpose
POST	/api/simulation/start	Start SUMO simulation
POST	/api/simulation/pause	Pause simulation
POST	/api/simulation/resume	Resume simulation
POST	/api/simulation/stop	Stop simulation
POST	/api/simulation/reset	Reset simulation
POST	/api/simulation/speed	Set speed multiplier
GET	/api/simulation/state	Current simulation state
GET	/api/vehicles	All vehicle states
GET	/api/signals	All signal states
POST	/api/signals/:id/state	Override signal state

Traffic Endpoints

Method	Path	Purpose
GET	/api/traffic	Live city traffic state
GET	/api/traffic/roads	Road registry
GET	/api/traffic/intersections	Intersection registry + live queues

Emergency Endpoints

Method	Path	Purpose
POST	/api/emergency	Create emergency event
GET	/api/emergency	List emergency events
GET	/api/emergency/:id	Emergency detail (route, position, ETAs)

Corridor Endpoints

Method	Path	Purpose
POST	/api/corridors	Create green corridor for emergency
GET	/api/corridors	List all corridors
GET	/api/corridors/:id	Corridor detail (windows, states)
POST	/api/corridors/:id/activate	Manually activate corridor
POST	/api/corridors/:id/cancel	Cancel corridor

Prediction & Analytics Endpoints

Method	Path	Purpose
GET	/api/predictions	Current junction predictions
GET	/api/analytics	Aggregated performance analytics

Scenario Comparison Endpoints

Method	Path	Purpose
POST	/api/scenarios/compare	Start baseline vs ITMS comparison
GET	/api/scenarios/compare/:jobId	Poll comparison job status
GET	/api/scenarios/runs	Recorded run metrics

Mobile API Endpoints

Method	Path	Purpose
POST	/api/mobile/auth/login	Driver authentication

GET	/api/mobile/emergency/active	Driver's active emergency
POST	/api/mobile/emergency/:id/evidence	Upload photo evidence
POST	/api/mobile/telemetry	Submit GPS telemetry

16. Real-Time Communication

WebSocket Endpoint

ws://127.0.0.1:3000/ws

Events (Backend → Frontend)

Event	Payload	Frequency
traffic:update	Full city traffic state (segments, junctions, summary)	~5 Hz
vehicle:update	All vehicle positions/speeds/angles	~5 Hz
signal:update	All signal states + queues	~5 Hz
prediction:update	Junction predictions (4 horizons each)	~2 Hz
emergency:created	New emergency event	On creation
emergency:update	Position, ETA, status update	~5 Hz
corridor:created	New corridor created	On creation
corridor:update	Corridor window states + status	~5 Hz
route:switched	Dynamic route change audit record	On switch
comparison:update	Job progress + delta results	On state change
system:alert	Disconnection / database error	On error
heartbeat	Liveness ping	Every 15s

Connection Behavior

- **Heartbeat:** Backend sends `heartbeat` every 15 seconds

- **Liveness check:** Frontend forces reconnect if silent > 35 seconds
- **Reconnect policy:** Exponential backoff (1s → 2s → 4s → ... → 30s cap)

17. Testing & Verification

Test Suite

Suite	Command	Tests	Status
Backend unit + integration	<code>npm run test:api</code>	132/132	✓ PASS
Python ML pipeline	<code>npm run test:prediction</code>	12/12	✓ PASS
TypeScript typecheck	<code>npm run typecheck</code>	All 3 packages	✓ PASS
Frontend ESLint	<code>npm run lint</code>	0 errors, 0 warnings	✓ PASS
Production build	<code>npm run build:web</code>	13 pages compiled	✓ PASS

Test Categories (Backend)

Category	Tests
TraCI codec	16
Traffic metrics	11
Network loader	5
Config validation	6
A* routing	12 (all 30 junction pairs)
Corridor planner + safety	17
Loop re-evaluation	4
TraCI client	7
Simulation manager	9

Vehicle reaction	1
Pipeline integration	3
API + WebSocket	4
Emergency integration	2
Prediction client	12
Prediction integration	4
Corridor integration	4
Closed-loop integration	7

18. Performance Metrics

Simulation Comparison Results

From the final E2E verification (ambulance W1→E2, 6-junction grid network):

Metric	Baseline	ITMS	Delta
Emergency Travel Time	162 s	82 s	-49.4%
Average Vehicle Delay	21.94 s	15.66 s	-6.28 s
Average Queue Length	25.59	21.78	-3.81
Average Speed	6.09 m/s	6.36 m/s	+0.27 m/s
Throughput	18.65/h	31.86/h	+70.8%

ML Inference Performance

Metric	Value
Single horizon (XGBoost)	1.01 ms
All 4 horizons	4.53 ms
Test R ² (30s horizon)	0.911

System Throughput

Metric	Value
WebSocket events/second	~14.7 events/s (5 Hz × all topics)
Simulation step interval	1000 ms (real-time)
Traffic collection	Lock-step with simulation
Prediction refresh interval	2000 ms

API latency, DB query time, and route calculation time not separately measured in current implementation.

19. Deployment Guide

Prerequisites

Requirement	Version	Installation
Node.js	≥ 24.7.0	nodejs.org
SUMO	1.27.1	<code>pip install eclipse-sumo</code>
PostgreSQL	16 + PostGIS 3.4	Docker Compose
Python	≥ 3.10	python.org
Docker Desktop	—	Optional (for DB)

Environment Variables

Backend (`apps/api/.env`):

```
PORT=3000

HOST=127.0.0.1

LOG_LEVEL=info

ITMS_DEMO=city

DEMO_CITY=Bhopal
```

```
DATABASE_URL=postgres://itms:change-me@127.0.0.1:5433/itms
```

```
TRAFFIC_EVENT_INTERVAL_MS=200
```

```
PREDICTION_SERVICE_URL=http://127.0.0.1:8100
```

```
VULTR_SERVERLESS_INFERENCE_API_KEY=your_key
```

```
VULTR_INFERENCE_MODEL=deepseek-v4.1-flash
```

Frontend (`apps/web/.env.local`):

```
NEXT_PUBLIC_API_URL=http://127.0.0.1:3000
```

Startup Sequence

```
# 1. Install dependencies
```

```
npm install
```

```
npm run build:network
```

```
npm run fetch:city && npm run build:city-network
```

```
docker compose --env-file docker/.env up -d db
```

```
cd services/prediction
```

```
pip install -r requirements.txt
```

```
python -m uvicorn inference.service:app --port 8100
```

```
npm run dev:api    # http://127.0.0.1:3000
```

```
npm run dev:web    # http://localhost:3001
```

Port Map

Service	Port
Backend API + WebSocket	3000
Frontend (dev)	3001
PostgreSQL (Docker)	5433
ML Prediction Service	8100

20. Limitations & Known Issues

Current Limitations

Limitation	Detail
Simulation only	ITMS operates on a SUMO digital twin, not real-world sensors
No production authentication	Development mode has no auth for the web UI
Single simulation instance	SUMO accepts only one TraCI connection; comparison runs are sequential
No pedestrian modeling	Sidewalks omitted from network; bounded-override stands in
ETA at intersection	Falls back to full-route ETA when vehicle is inside an intersection (internal edge)
Spawn delay	Emergency vehicle spawn may be deferred by SUMO if departure edge is blocked

Model quality	XGBoost trained on 6-junction synthetic network; +120s horizon $R^2 = 0.74$
Comparison run time	~30–60s wall clock per run at full simulation speed
Browser E2E	Playwright click-through not fully verified in documentation environment

Fixed Issues (Historical)

- **CORS:** Missing `@fastify/cors` caused 404 preflights. Fixed with proper array/boolean origin support.
- **Tailwind v4:** Globals.css used v3 directives causing zero utilities. Fixed with `@import "tailwindcss"`.
- **MapLibre:** Style spec error and Next.js worker URL issue. Fixed by complete removal + SVG renderer.
- **Google Maps:** Later removed per architecture decision (SUMO is the authoritative map).
- **Comparison runner:** Failed when simulation was already running. Fixed: stops existing simulation first.
- **TraCI VAR_ANGLE:** Variable ID was incorrect (0x0d instead of 0x43). Fixed and verified against SUMO source.

21. Implementation Status Matrix

Component	Status	Evidence
SUMO 6-junction grid network	✅ IMPLEMENTED + VERIFIED	Network built, 132 tests passing
SUMO Bhopal city network	✅ IMPLEMENTED + VERIFIED	883 segments, georeferenced
TraCI TCP client	✅ IMPLEMENTED + VERIFIED	16 codec tests + 7 TraCI tests
Traffic collector (per-step)	✅ IMPLEMENTED + VERIFIED	11 metrics tests + pipeline tests
PostgreSQL persistence	✅ IMPLEMENTED + VERIFIED	6 migrations, run-scoped records

PostGIS spatial storage	✅ IMPLEMENTED + VERIFIED	SRID 0, geometry stored and served
WebSocket real-time events	✅ IMPLEMENTED + VERIFIED	~14.7 events/s verified live
A* routing engine	✅ IMPLEMENTED + VERIFIED	All 30 grid pairs routable; tested
Emergency vehicle creation	✅ IMPLEMENTED + VERIFIED	API creates, SUMO spawns, tracked
Per-junction ETA calculation	✅ IMPLEMENTED + VERIFIED	Live ETAs in API response
Arrival detection	✅ IMPLEMENTED + VERIFIED	SUMO arrived-vehicle list
XGBoost ML service	✅ IMPLEMENTED + VERIFIED	R ² 0.911 (30s), 132 tests
Prediction API (FastAPI)	✅ IMPLEMENTED + VERIFIED	POST /predict, GET /health
Prediction fallback	✅ IMPLEMENTED + VERIFIED	Rule-9: recent valid → deterministic
Green corridor planner	✅ IMPLEMENTED + VERIFIED	17 planner/safety tests
Safety constraint engine	✅ IMPLEMENTED + VERIFIED	8 checks, conflict-free by construction
Corridor executor (rolling)	✅ IMPLEMENTED + VERIFIED	Apply/clearance/restore verified live
Signal program restoration	✅ IMPLEMENTED + VERIFIED	TraCI setProgram on passage/cancel
Closed-loop optimization	✅ IMPLEMENTED + VERIFIED	7 closed-loop integration tests
Hysteresis-guarded rerouting	✅ IMPLEMENTED + VERIFIED	Route switch audit persisted
Scenario comparison runner	✅ IMPLEMENTED + VERIFIED	Baseline 162s vs ITMS 82s measured
Simulation metrics recorder	✅ IMPLEMENTED + VERIFIED	Per-run finalization to DB

Command Center UI (9 pages)	✅ IMPLEMENTED + VERIFIED	Production build ✓, SSR verified
SVG simulation map	✅ IMPLEMENTED + VERIFIED	Real SUMO geometry, real TraCI coords
Mobile API (drivers, fleet)	✅ IMPLEMENTED	Code and DB schema exist
AI Copilot (Vultr/LLM)	✅ IMPLEMENTED	Requires VULTR_API_KEY to activate
Roadside IoT displays	✅ IMPLEMENTED	Code and DB schema exist
Playwright browser E2E	⚠️ NOT VERIFIED	Planned as next step